



Input



LLM judge f
model under audit



Item x
one pairwise
comparison



Reference labels
calibration / test



Risk budget
alpha, delta
target selective risk

CARE-Judge

from protocol perturbations to a risk-controlled verdict



Protocol perturbations

6 judge calls per item



base
verdict
 $\tau=0$



position
swap



rubric
paraphrases
 $K=3$



self-
consistency
 $S=3$

per-item
traces



Feature families $\phi(x)$



rubric stability



position stability



self-consistency



confidence, length

calibrated
scores



Correctness calibrator

$$\hat{p}(\text{correct} \mid \phi(x))$$

fit on **TRAIN**



Risk-controlled threshold

$\hat{\lambda}$: exact Clopper-Pearson

selected on **CAL**

Selective verdict

on held-out **TEST**



ACCEPT

$$\hat{p} \geq \hat{\lambda}$$

certified verdict



ABSTAIN

$$\hat{p} < \hat{\lambda}$$

risky item or a judge
too weak to certify



Reported:

coverage,
selective risk,
violation rate